

# How Computers Represent Numbers

John Finigan

November 20, 2015

# Typical modern computer hardware

- ▶ PC, iPhone, IBM Mainframe
- ▶ All information stored in memory as base-2 digits
- ▶ Every byte (8 base-2 digits) has a unique address
  - ▶ A byte is our fundamental unit of information

# Bytes and Words

- ▶ 1 byte is 8 bits (8 base-2 digits)
- ▶ 1 byte can store  $2^8 = 256$  unique values
- ▶ So, we group bytes into words
- ▶ 4 byte word can store  $2^{32} = 4,294,967,296$  unique values
- ▶ 2, 4, and 8 byte words are common sizes for numbers

# Representing $\mathbb{Z}$ - The Integers

- ▶ Consider a word as one base-2 number
- ▶ A 2 byte word is then one 16 digit base-2 number
- ▶ Range:  $0000000000000000_2$  to  $1111111111111111_2$
- ▶ In base-10, that's  $0_{10}$  to  $65,535_{10}$

# First problem: Limited Range

- ▶  $\mathbb{Z}$  is infinitely large
- ▶ No getting around this mismatch
- ▶ Minimum value: 0
- ▶ Maximum values:
  - ▶ 2 byte word -  $65,535_{10}$
  - ▶ 4 byte word -  $4,294,967,295_{10}$
  - ▶ 8 byte word -  $18,446,744,073,709,551,615_{10}$

## Second problem: No Negative Numbers

- ▶ Solution: Reserve half of the range for negative numbers
- ▶ Zero is placed in the center of our finite number line
- ▶ Largest absolute value is cut in half
- ▶ One bit is consumed by the sign, one way or the other

# Representing Negative Numbers, Take 1

- ▶ Sign/Magnitude representation
  - ▶ Most significant bit indicates sign
  - ▶ Rest indicate magnitude (absolute value)
  - ▶ Example in 4 bits:  $0111_2 = 7_{10}$  negates to  $1111_2 = -7_{10}$
- ▶ Ones Complement
  - ▶ Invert every bit in a positive number to make a negative number with same absolute value
  - ▶ Example in 4 bits:  $0111_2 = 7_{10}$  negates to  $1000_2 = -7_{10}$

# Problem! Negative Zero

- ▶ Not impossible to work around, just awkward
- ▶ 4 bit examples:
  - ▶ Sign/Magnitude:  $0000_2$  and  $1000_2$
  - ▶ Ones Complement:  $0000_2$  and  $1111_2$
- ▶ To test equality with zero, must test both
  - ▶ Either in hardware, or in software



# Representing Negative Numbers, Take 2

- ▶ Twos Complement
  - ▶ To make a negative number, invert every bit in its positive counterpart and then add 1
  - ▶ Example in 4 bits:  $0111_2 = 7_{10}$
  - ▶ Negates to  $1000_2 + 1_2 = 1001_2 = -7_{10}$
- ▶ No Negative Zero:
  - ▶  $0000_2$  inverts to  $1111_2$
  - ▶  $1111_2 + 1_2 = 0000_2$ , discarding carry
  - ▶ Twos complement negative of  $0000_2$  is  $0000_2$
- ▶ Adding twos complement numbers “works”:
  - ▶  $7_{10} + -7_{10} = 0$
  - ▶  $0111_2 + 1001_2 = 0000_2$ , discarding carry

# Overflow

- ▶ Notice the “discarding carry” from previous slide
- ▶ A word can’t grow, so carry is lost
- ▶ In other words, modular arithmetic
- ▶ For a two byte word,  $65,535_{10} + 1_{10} = 0$
- ▶ C guarantees modular arithmetic for unsigned integers, but guarantees nothing for signed integers: result is undefined.

# Unsigned Overflow demo - C language

```
unsigned int u = UINT_MAX;  
printf("UINT_MAX is %u\n", u);  
printf("UINT_MAX + 1 is %u\n", u + 1);  
printf("UINT_MAX + 2 is %u\n", u + 2);
```

Output:

```
UINT_MAX is 4294967295  
UINT_MAX + 1 is 0  
UINT_MAX + 2 is 1
```

Signed overflow in C is undefined!

# Representing $\mathbb{R}$ - The Reals

- ▶ Range still limited by finite word size
- ▶ New problem: real number line is infinitely dense
  - ▶ Base-2 word is fundamentally an integer
  - ▶ Unavoidably, it can represent only a finite set of values
- ▶ Not all real values can have unique/exact representations

# Floating Point

- ▶ Divide the word into two major parts
- ▶ One part stores a binary value with a fixed radix point
  - ▶ Called the Significand or Fraction or Mantissa
  - ▶ Example:  $1.01010101_2$
- ▶ The other part stores a binary value representing an integer exponent
  - ▶ An implicit base,  $\beta$ , is raised to this exponent
- ▶ Multiplying the significand by the exponentiated base yields the word's real number value
  - ▶ Value is computed as  $\pm d.dddddddd \times \beta^e$

# Exact representation of a real number is possible

- ▶ But only for certain reals:
  - ▶ If the desired value is not too large or too small
  - ▶ If the desired value and the base share the same prime factors
    - ▶ Base-10 example:  $0.1 = \frac{1}{10} = \frac{1}{2 \cdot 5}$
    - ▶ If  $\beta = 2$ , 5 is not a common prime factor!
- ▶ Uniqueness isn't guaranteed either:
  - ▶  $1.0 \times 10^1 = 0.1 \times 10^2$
  - ▶ But, if we require that the significand  $\geq 1.0$ , uniqueness is guaranteed
  - ▶ A floating point number with this property is *normalized*.

# Example Encoding

- ▶ 4 byte word:

- ▶ 1 bit of sign, 8 bits of exponent, 23 bits of fraction
- ▶ S **EEEE EEEE** **FFFF FFFF FFFF FFFF FFFF FFFF**

- ▶ 8 byte word:

- ▶ 1 bit of sign, 11 bits of exponent, 52 bits of fraction
- ▶ S **EEEE EEEE EEEE** **FFFF FFFF FFFF FFFF FFFF FFFF FFFF**
- ▶ **FFFF FFFF FFFF FFFF FFFF FFFF FFFF FFFF**

# Encoding Choices

- ▶ How to represent sign of fraction and sign of exponent
- ▶ What base,  $\beta$ , is raised to the exponent?
- ▶ How to represent special values, like 0, infinity, or undefined
- ▶ Many competing obsolete defacto standards
  - ▶ IBM Mainframe HFP,  $\beta = 16$
  - ▶ DEC VAX
  - ▶ Cray
  - ▶ dozens more



# IEEE 754 Standard (Basic Formats)

- ▶ 4 byte word (single precision):
  - ▶ 1 bit of sign, 8 bits of exponent, 23 bits of fraction
- ▶ 8 byte word (double precision):
  - ▶ 1 bit of sign, 11 bits of exponent, 52 bits of fraction
- ▶ For each:
  - ▶  $\beta = 2$
  - ▶ Sign/Magnitude sign representation for fraction
  - ▶ *biased* sign representation for exponent
  - ▶ Exponent's unsigned value is added to a constant *bias*

# IEEE 754 Limits

- ▶ Single precision

- ▶  $e_{max} = 127, e_{min} = -126$
- ▶ 23 bits of fraction plus implicit most significant bit (1.)
- ▶ Normalized range: about  $\pm 1.175 \times 10^{-38}$  to  $\pm 3.403 \times 10^{38}$

- ▶ Double precision

- ▶  $e_{max} = 1023, e_{min} = -1022$
- ▶ 53 bits of fraction plus implicit most significant bit (1.)
- ▶ Normalized range: about  $\pm 2.225 \times 10^{-308}$  to  $\pm 1.798 \times 10^{308}$

# IEEE 754 Special Values (Single Precision)

- ▶ If raw exponent is 255 (all 1s) and fraction is not 0
  - ▶ NaN - Not a Number, sign irrelevant
- ▶ If raw exponent is 255 (all 1s) and fraction is 0
  - ▶ Infinity - sign bit determines sign
- ▶ If raw exponent is 0 (all 0s) and fraction is not 0
  - ▶ Denormalized numbers, sign bit determines sign
- ▶ If raw exponent is 0 (all 0s) and fraction is 0
  - ▶ Zero - sign bit determines sign

# Error

- ▶ “Squeezing infinitely many real numbers into a finite number of bits requires an approximate representation.” - Goldberg
- ▶ Units in the last place (ulps)
  - ▶ How much does the rounded result differ from the ideal result?
  - ▶ 3.14 approximated to 3.12 - 2 ulps
  - ▶ 3.14159 approximated to 3.14 - 0.159 ulps
  - ▶ Due to rounding, approximation can differ from ideal by up to 0.5 ulps
- ▶ Relative error:
  - ▶  $(\text{ideal} - \text{approx}) / \text{ideal}$
  - ▶ Relative error corresponding to 0.5 ulps varies by at most a factor of  $\beta$  (Goldberg)

# Overflow and Underflow

- ▶ IEEE 754 mandates that overflow and underflow can be detected
  - ▶ a flag is set and, optionally, an exception can be raised
- ▶ Gradual Underflow vs. Store-0
  - ▶ Recall that the fraction's most significant bit is implicitly a 1 before the radix point
  - ▶ This leaves a BIG gap of unrepresentable numbers between  $2^{e_{min}}$  and zero
  - ▶ Old non-IEEE 754 implementations typically store 0 for these numbers

# Gradual Underflow

- ▶ IEEE 754 introduces a special case: if exponent is  $e_{min} - 1$  and fraction is non-zero, implicit digit before radix point is 0.
  - ▶ called a subnormal or denormalized number
- ▶ We can then use all of the fraction's bits to represent values between  $2^{e_{min}}$  and 0
- ▶ Calculations do not have to be scaled to avoid tiny values that, with Store-0, would be represented *less accurately* than normalized numbers
- ▶ Makes error analysis and avoidance easier
  - ▶ Otherwise, algorithms have to detect when values get too small, and if possible, scale to avoid them

## Error Example

```
double a = 1.015;  
printf("%19.16f\n", a);
```

Output

1.0149999999999999

$1.015_{10}$  is not exactly representable in base-2 floating point.

# Error can break your program

```
//Goldberg's constants  
double x = (3.34*3.34) - (4.0*1.22*2.28);  
printf("%19.16f\n", x);  
  
if (x == 0.0292)  
    printf("equal to 0.0292!\n");  
else  
    printf("not equal to 0.0292!\n");
```

Output:

```
0.02920000000000012  
not equal to 0.0292!
```



# Sources of Error beyond 0.5 ulps

- ▶ Error accumulates
- ▶ Expressions may not be associative or distributive
- ▶ Summations can accumulate error
  - ▶ Using naive approach, can grow proportional to number of terms summed
  - ▶ My own experiment: relative error of summing 1.015
  - ▶ Summed 100 times vs 1 trillion times: relative error 29 billion times larger!
- ▶ Cancellation: subtracting two numbers that are almost equal can cancel most of the correct digits, leaving mostly incorrect digits.

# Solution

It's an entire field of study!

- ▶ Many recipes have been developed, and many proofs done
- ▶ Active field since at least the 1960s
- ▶ Informed IEEE 754 deeply
- ▶ Many “simple” equations can be rearranged in potentially complex ways to avoid error.
- ▶ Example: Kahan's summation algorithm reduces error growth to a constant factor

# Decimal Floating Point

- ▶ IEEE 854
- ▶  $\beta = 10$
- ▶ We (mostly) think in base-10, so it's better for human generated numbers
- ▶ Such as money
- ▶ Same underlying approximation issues though
- ▶  $\beta > 2$  can make error analysis more difficult (Goldberg), but 10 is a special case for human reasons

# Software defined precision

- ▶ IEEE 754 happens in hardware (typically)
- ▶ We can do higher precision more slowly using software
- ▶ GNU Multiple Precision Arithmetic Library (GMP)
  - ▶ " There is no practical limit to the precision except the ones implied by the available memory in the machine GMP runs on."

# References

- ▶ David Goldberg. 1991. What every computer scientist should know about floating-point arithmetic. ACM Comput. Surv. 23, 1 (March 1991)
- ▶ IEEE STANDARD 754-1985 - IEEE Standard for Binary Floating-Point Arithmetic
- ▶ Sun Numerical Computation Guide (2005, Sun Microsystems inc.)